

BEYOND RETENTION: ACTIVE FORGETTING IN MULTI-TURN LLM AGENTS

Anonymous authors

Paper under double-blind review

ABSTRACT

Multi-turn LLM degradation on underspecified conversations is largely an artifact of the chat protocol, not a model-capability deficit. On Lost-in-Conversation with Qwen2.5-14B, a five-line policy that drops the model’s own prior tokens at the final turn (FRESH-LAST) recovers 87% of the multi-turn gap; dropping them at every turn (FRESH-EVERY) reaches 103% and exceeds the single-turn ceiling. The fix is binary rather than graded: keeping any 25–75% of prior assistant tokens recovers only 4–23% of the gap. The mechanism is attention-mediated. On $n=22$ failed sharded conversations (Qwen2.5-7B), prior provisional commitments receive per-token attention comparable to the system prompt and 8.5% above user evidence, so attention-salience-based KV compression preferentially keeps tokens that should be evicted. Text-level mitigations also fail: prefixing prior tokens with [OUTDATED] barely moves accuracy. The effect replicates across five models in four architecture families (Qwen-7B/14B, Llama-3.1-8B, Mistral-7B, Gemma-2-9B). We frame this as a missing memory primitive. Agent systems implement *generation* and *retention* (KV cache reuse) but lack *active forgetting*. We introduce VCHR and HCHR alongside the standard prefix-cache hit rate (PCHR; message-level proxy mPCHR) to make forgetting measurable, and propose CACHEGUARD, a typed forgetting controller that drops provisional commitments while retaining verified evidence. CACHEGUARD validates on all four architecture families (Qwen-14B 0.875, Llama-3.1-8B 0.704, Mistral-7B 0.544, Gemma-2-9B 0.741), and matches or exceeds FRESH-EVERY on Mistral and Gemma. The picture reverses cleanly on stateful tool use (BFCL v4): FRESH-EVERY collapses to 0.090 while CACHEGUARD matches the strongest baseline (0.280 vs. 0.295, $p=0.66$, ns), so forgetting must be content-typed rather than uniform. Chat protocols and serving stacks should expose active forgetting as a first-class primitive, alongside generation and retention.

1 INTRODUCTION

LLM agent systems are designed around two memory operations: *generation* (the agent’s outputs) and *retention* (KV-cache reuse, optimized in modern serving stacks like vLLM (Kwon et al., 2023), SGLang (Zheng et al., 2024), and LMCache (Liu, 2025), monitored via prefix-cache hit statistics). A third operation has received much less attention: *active forgetting*. Standard chat protocols have no forgetting mechanism beyond full-history truncation: a user message (external evidence) and a prior model output (often a provisional commitment under partial information) both persist indefinitely, serialized under identical role tags. Existing serving systems are intentionally agnostic to the epistemic status of cached content—they answer “can this prefix be reused?” but not “should this content remain eligible for future decisions?”. The latter question is what forgetting answers.

A multi-turn agent context is not a static document. A user message is external evidence. A successful tool observation is external evidence. But an assistant response is a model action—often a provisional commitment made under partial information—which the standard chat protocol then serializes into the next prompt as if it were evidence. Recent work (Laban et al., 2025) reports that multi-turn LLM accuracy on underspecified conversations drops 30–40% relative to single-turn fully-specified versions, even though by conversation’s end the assistant has received the same information. Existing remedies—prompt evolution (Agrawal, 2025), harness search (Lee et al., 2026; Zhang, 2025), dialogue policy learning (Chen, 2025), and the LiC paper’s own RECAP—all

retain the assistant’s prior responses in context. RECAP attains the highest message-level prefix-cacheability proxy (mPCHR 0.851 in our measurement) yet recovers only 65% of the multi-turn gap, 22 pp below FRESH-LAST. Optimizing for cacheability has a hidden cost.

We observe that the LLM’s prior responses, generated under partial information when the user has not yet revealed all shards, are systematically wrong but stored in context with identical chat-template formatting as user evidence. They are physically reusable (every prefix-cache lookup retrieves them) but epistemically invalid (the conversation’s later user shards now contradict them). We propose the distinction: *can reuse* $\neq$ *should reuse*. To make this operational, we introduce two metrics alongside PCHR:

- **VCHR** (valid cache hit rate): of reused tokens, the fraction belonging to blocks that remain epistemically valid evidence under the current state.
- **HCHR** (harmful cache hit rate): of reused tokens, the fraction whose presence flips the answer from correct to wrong (counterfactually).

On Lost-in-Conversation with Qwen2.5-14B, the standard SHARDED protocol (no forgetting) attains mPCHR 0.751 but accuracy 0.208. RECAP, which re-states user information but does not forget assistant outputs, attains the highest mPCHR (0.851) and accuracy 0.625. A five-line forgetting policy that drops all prior assistant tokens at every turn (FRESH-EVERY) attains mPCHR 1.000, VCHR 1.000, and accuracy 0.875. The 22 pp gap between RECAP and FRESH-LAST (65% vs 87% recovery, $\Delta = +0.146$) isolates to the single line distinguishing them: whether the prior assistant trace is forgotten at decision time. Four findings sharpen this picture:

- **Cliff, not gradient.** Dropping 25–75% of prior assistant messages recovers only 4–23% of the gap. Only complete removal ($K = 1.0$, FRESH-LAST) recovers 87%; no partial compromise works.
- **Capability inversion.** Qwen2.5-7B under FRESH-LAST reaches 0.917, nominally above 14B FRESH-LAST (0.771) and the 14B Concat ceiling (0.854). Larger models generate more coherent commitments under partial information, producing stronger anchors rather than resistance. The nominal exceedance is partially attributable to multi-attempt advantage on Qwen-7B (Appendix C); on Llama-3.1-8B the same ordering survives a best-of- N control.
- **Attention, not format.** Marking prior messages with [OUTDATED PARTIAL-INFO COMMITMENT] yields 0.307, barely above the 0.208 baseline and far below FRESH-LAST (0.771). Text markers do not override attention-mediated anchoring.
- **Task-structure dependent.** RECAP recovers 86% of the gap on tool-use actions but only 27% on code generation (Table 1); single-call commitments are easier to partially override than multi-step ones.

We synthesize these findings into CACHEGUARD, a typed forgetting controller. It (i) classifies each context block by epistemic role—system instruction, user evidence, tool result, assistant artifact, assistant claim, assistant plan; (ii) applies an eligibility rule that retains user/tool blocks but forgets provisional assistant claims; and (iii) renders a stable-valid prefix plus a volatile suffix at LLM-call time. CACHEGUARD is deployable client-side with no training, no model change, and no proprietary serving access.

Our contributions are:

- **Framing.** We position *active forgetting* as a missing primitive in LLM agent memory, alongside generation and retention. Three metrics (PCHR, VCHR, HCHR) operationalize it: PCHR measures retention, VCHR measures whether reused content is still valid evidence, HCHR measures whether reused content is causally harmful.
- **Diagnosis.** RECAP and FRESH-LAST differ only in whether prior assistant tokens are forgotten at decision time, yet FRESH-LAST recovers 22 pp more of the gap. The dose-response is a cliff at $K=1.0$, a ceiling exceedance survives a best-of- N control on Llama-3.1-8B (Appendix C), and [OUTDATED] markers leave accuracy at 0.307 vs. 0.771 for full removal. Forgetting is binary, attention-mediated, and not overridable by text instructions.
- **Method.** CACHEGUARD classifies context blocks by epistemic role, drops provisional commitments, and retains verified evidence. It needs no retraining and is deployable client-

side. The forgetting effect and CACHEGUARD both replicate across five models in four architecture families (Qwen-7B/14B, Llama-3.1-8B, Mistral-7B, Gemma-2-9B); on Mistral and Gemma CACHEGUARD matches or exceeds FRESH-EVERY, showing that typed eligibility can outperform uniform forgetting.

- **Two-regime validation.** On LiC (stateless reasoning) FRESH-EVERY reaches 0.875. On BFCL v4 (stateful tool use) FRESH-EVERY collapses to 0.090 (−20.5 pp vs. SHARDED, $p < 0.001$, $N = 200$). CACHEGUARD matches SHARDED on BFCL (ns) while still beating it on LiC, so forgetting must be content-typed rather than uniform.

2 RELATED WORK

Multi-turn LLM degradation and mitigation. Laban et al. (2025) document the multi-turn underspecified failure mode and propose RECAP and SNOWBALL, which re-state user information but retain prior assistant responses, recovering 30–52% of the gap. The same study shows that reasoning models (o3, DeepSeek-R1) degrade similarly: extra test-time compute does not, on its own, help models strategize across turns. Weston (2025) corroborate this on MultiChallenge, where all frontier models score under 50% on instruction retention and self-coherence. Li (2026b) propose a Mediator-Assistant architecture that rewrites user inputs while again preserving assistant history. Harness and prompt search systems (Lee et al., 2026; Zhang, 2025; Agrawal, 2025) optimize scaffolding or prompt text but operate within a fixed context-retention protocol. The dialogue-systems community addressed multi-turn state maintenance via explicit slot-filling (Wu et al., 2019); our finding shows that implicit state tracking through LLM history retention is insufficient, and selective forgetting of agent commitments outperforms retention on knowledge-intensive tasks. Our diagnosis shows that a protocol-level change dominates all prompt-level interventions on LiC.

Closest concurrent work: assistant-history filtering. Huang et al. (2026) report that omitting prior assistant responses preserves quality and reduces context up to $10\times$ on in-the-wild conversations, identifying “context pollution” on 36.4% of self-contained turns. We extend this in four ways: (i) controlled sharded benchmark where the failure is reproducible by construction; (ii) a binary-cliff mechanism showing partial filtering does not work; (iii) formalization as a *cache validity* problem with three metrics connecting to the serving stack; (iv) CACHEGUARD, a typed reuse layer that handles the stateful-tool reversal (Section 5.7) where assistant-history deletion is harmful. Three concurrent papers identify the same anchoring phenomenon under different names. Wang & Sun (2025) introduce PI-LLM, showing log-linear decay in retrieval accuracy under proactive interference even when the target value is adjacent to the query; their key result—“ignore earlier inputs” instructions fail—directly corroborates our marked-history null. Li (2026a) call the effect “conversational inertia” and propose periodic blind context clipping; Zhang (2026) address it via expensive RL fine-tuning. Xu (2026) build a learned sleep-cycle gate inside the KV cache. Compared with these, our contribution is the cliff mechanism, the typed-vs-untyped reversal on stateful tool use, and a zero-shot, training-free, deployable controller.

Active forgetting in cognitive architectures and agent memory. Forgetting is a first-class operation in cognitive architectures: directed forgetting (Bjork, 1972) and ACT-R activation decay (Anderson & Lebière, 1998) explicitly evict working-memory contents. Anchoring bias in human judgment (Tversky & Kahneman, 1974)—where numerical estimates remain biased toward an initial value even when that anchor is identified as arbitrary—provides a behavioral analog to commitment contamination: our LLM-level effect differs in that the anchor is not arbitrary but is the model’s own prior response, making it a self-reinforcing rather than externally imposed bias. Recent LLM agent systems revive this lineage (Park et al., 2023; Packer et al., 2023; Li et al., 2025b; Kosinski, 2025; ByteDance-Tsinghua-SIA, 2026) with hierarchical memories and importance-based eviction, but eviction is treated as a pragmatic optimization rather than a correctness-bearing primitive. Our framing sharpens this for chat-protocol multi-turn settings: forgetting is content-typed (provisional commitment vs. verified evidence) and is structurally necessary, not optional. We distinguish active forgetting at *inference time* on the KV cache from machine unlearning, which operates on *training-time weights*; these address categorically different failure modes (episodic-context bias vs. parametric-knowledge erasure). Khattab et al. (2025) survey context engineering; our work fills the missing typed-eviction module in their taxonomy.

KV cache systems, position bias, and compression. Kwon et al. (2023), Zheng et al. (2024), and Liu (2025) optimize physical cache hit rate. Zhang et al. (2023); Xiao et al. (2024); Li (2024) compress the KV cache by attention salience. All treat reusable as reuse-worthy and select tokens to retain. The over-attended tokens in our setting are content-anchored prior assistant commitments, distinguishing them from position-zero attention sinks (Xiao et al., 2024; Sun et al., 2025) and from the U-shaped position bias of Liu et al. (2024). The cliff dose-response (Section 5.3) rules out salience-based compression as a fix—any retention heuristic that keeps a non-trivial fraction of prior assistant tokens fails. Anonymous (2026) document a related token-level anchoring effect; we quantify it through cache-validity metrics and provide a deployable forgetting policy. Recursive summarization (Wu et al., 2023) compresses but retains assistant-generated content, which our marked-history null predicts will inherit the anchoring effect.

Sentence-level anchoring and causal attribution. Li et al. (2025a) identify “thought anchors” within chain-of-thought reasoning: individual sentences whose removal flips a wrong-to-correct prediction, finding a small fraction ($\sim 5\%$) of CoT sentences carry disproportionate causal weight. Their scope is *intra-turn*—sentences within a single response. Our finding identifies the analogous phenomenon *cross-turn*: entire prior assistant responses from earlier conversation turns act as anchors on the final response, via the same attention-mediated mechanism (Section 5.5). Both are causal (removal flips outcome) and localizable via ablation; Appendix F extends their sentence-level methodology to cross-turn assistant history. Chen et al. (2026) show that LLMs in extended agentic conversations develop “AI psychosis”—persistent incorrect beliefs resisting correction under explicit counter-evidence—consistent with our contamination framing and the Marked-History null. Causal attribution methods (Kim et al., 2025) trace output causation to input tokens; HCHR operationalizes the same question at the block level, connecting it to the serving-layer design of Section 4.3.

3 BACKGROUND AND PROBLEM SETTING

Sharded multi-turn underspecification. We use the Lost-in-Conversation *sharded* benchmark (Laban et al., 2025), which transforms a single-turn fully-specified instruction into a sequence of K *shards*—atomic information units, each revealed in a separate user turn. Across K turns, the user simulator reveals one shard per turn until all are revealed. The assistant produces a response after each user turn. The conversation terminates either when an automatic verifier judges an assistant response as a correct answer attempt, or when all shards have been revealed. We evaluate on the LiC `multi_heldout` split covering three task families with non-trivial Concat ceilings: math (GSM8K), code (HumanEval / LiveCodeBench), and tool-use (BFCL actions).

Reference protocols. LiC defines two reference protocols. CONCAT concatenates all shards in a single user message and queries the LLM once; this is the single-turn ceiling. SHARDED runs the protocol described above and is the floor. The gap between these defines the multi-turn drop. LiC also reports two mitigations: SNOWBALL re-states all prior shards at each turn (keeps full history), and RECAP appends a final user message summarizing all revealed shards (also keeps full history). Both partially close the gap.

4 METHOD

4.1 THREE CACHE METRICS FOR MULTI-TURN LLM SERVING

Standard serving systems report *physical cache hit rate*:

$$\text{PCHR} = \frac{\text{reused prefix tokens}}{\text{queried prompt tokens}}. \quad (1)$$

We complement it with two metrics that depend on the conversation’s evolving epistemic state. Let $\mathcal{B}$ be the set of context blocks reused at decision time, and let $\text{valid}(b, s)$ denote whether block b is still epistemically valid evidence under current state s (Section 4.3).

$$\text{VCHR} = \frac{\sum_{b \in \mathcal{B}} \mathbf{1}[\text{valid}(b, s)] \cdot |b|}{\sum_{b \in \mathcal{B}} |b|}. \quad (2)$$

$$\text{HCHR} = \frac{\sum_{b \in \mathcal{B}} \mathbf{1}[\text{flips}(b)] \cdot |b|}{\sum_{b \in \mathcal{B}} |b|}, \quad (3)$$

where $\text{flips}(b)$ is 1 iff removing b counterfactually changes the outcome from incorrect to correct. PCHR alone defines the standard objective; an ideal serving stack maximizes $\text{PCHR} \times \text{VCHR}$ while minimizing HCHR. Existing systems measure only PCHR; multi-turn agentic deployments require also tracking VCHR and HCHR.

4.2 PROTOCOL VARIANTS

All protocols differ only in which subset of $\{a_1, \dots, a_{t-1}\}$ is passed to the LLM at turn t . **SHARDED**: full history $[s, u_1, a_1, \dots, u_t]$. **FRESH-LAST**: at the final turn, pass $[s, u_{\text{concat}}]$ dropping all prior assistant messages (five lines of code). **FRESH-EVERY**: at every turn t , pass $[s, u_{\text{concat-so-far}}]$; the LLM never sees its own prior responses. **MARKED-HISTORY**: retain all prior assistant messages but prefix each with `[OUTDATED PARTIAL-INFO COMMITMENT]`, testing whether anchoring is format-confusion (H3) or attention-mediated (H1). **SELECTIVE-DROP**: parametrized by $K \in [0, 1]$, drops the first (or last) K fraction of prior assistant messages; $K = 0$ is Sharded, $K = 1$ is Fresh-Last.

4.3 CACHEGUARD: VALIDITY-AWARE REUSE

The protocols above reveal a Pareto problem: **SHARDED** and **RECAP** have high mPCHR but low VCHR; **FRESH-EVERY** has high VCHR but its full-context-rebuild approach forfeits potential cache reuse on a stable system+constraint prefix. A practical serving layer would combine high PCHR (cacheable stable prefix) with high VCHR (no provisional commitments retained). This is what **CACHEGUARD** provides.

Block Parser. We type each non-log message in the conversation:

- `system` — protocol instruction (always valid).
- `user_evidence` — user-revealed shard (always valid).
- `tool_observation` — successful tool result (always valid).
- `assistant_artifact` — a code block, table, or named output produced by the assistant.
- `assistant_claim` — an answer attempt or boxed/bracketed final answer.
- `assistant_clarification` — a question to the user.
- `assistant_plan` — a description of intended next steps without producing artifact.

We additionally track `contradicted_by` (blocks invalidated by a later block) and `referenced_by` (artifacts referenced by later user or tool blocks).

Eligibility Controller. A block b is reuse-eligible iff:

- it is system, user, or successful tool; or
- it is an artifact and at least one later user/tool block references it; or
- it is a clarification (asking, not committing).

Provisional claims and plans are excluded by default. Contradicted blocks are excluded.

Renderer. At LLM-call time, the rendered context is

$$\underbrace{[s \parallel \text{verified user constraints} \parallel \text{verified tool state}]}_{\text{stable valid prefix}} \parallel \underbrace{[\text{current turn} \parallel \text{referenced artifacts}]}_{\text{volatile suffix}}.$$

The stable valid prefix is monotone-growing (only appended-to as new verified evidence arrives), maximizing PCHR; the volatile suffix is small and changes per turn. Provisional commitments live in a hidden side-buffer that the renderer does not pass to the LLM.

Table 1: Per-class accuracy and gap recovery on LiC `multi_heldout` (Qwen2.5-14B). Sharded is the multi-turn floor; Concat is the single-turn ceiling. Gap recovery = $(\text{acc} - \text{Sharded}) / (\text{Concat} - \text{Sharded})$ averaged across the three classes.

Protocol	Math	Code	Actions	Gap recovered
Sharded (floor)	0.500	0.000	0.125	0%
Concat (ceiling)	0.938	0.625	1.000	100%
Recap (Laban et al., 2025)	0.833	0.167	0.875	65%
Fresh-Last (ours)	0.875	0.438	1.000	87%
Fresh-Every (ours)	0.938	0.688	1.000	103%

Renderer detail: user-shard consolidation. The renderer concatenates verified user shards into a single block in the stable prefix, rather than preserving multiple user-role messages. This is consistent with the typed-eligibility view—the user’s accumulated evidence is a single epistemic unit—and matches how the task’s underlying single-turn evaluator (`populate_concat_prompt`) presents the same content. On LiC the consolidation is exact (each shard contributes one bullet); on BFCL where the user’s turn structure conveys timing, we keep user messages atomic. The choice is part of the typed renderer, not a benchmark hack.

5 EXPERIMENTS

5.1 SETUP

We use Qwen2.5-14B-Instruct as the assistant model and Qwen2.5-7B-Instruct as the user simulator and grader. We evaluate on the LiC `multi_heldout` split: 8 task instances per class across 3 task classes (math, code, actions) for 24 instances, each run with $N = 2$ independent simulations. Per-class accuracy is computed by averaging per-instance accuracy (treating null/no-answer outcomes as failures), then averaging across instances within each class. The gap is measured as $\text{Concat} - \text{Sharded}$; gap recovery for a method M is $(M - \text{Sharded}) / (\text{Concat} - \text{Sharded})$.

5.2 MAIN RESULT: PROTOCOL COMPARISON

Table 1 reports per-class accuracy and gap recovery. **Fresh-Last recovers 87% of the gap, 22 pp above LiC’s Recap baseline (65%). Fresh-Every recovers 103%, exceeding the single-turn ceiling on average.**

5.3 MECHANISM: ANCHORING IS A CLIFF-DOMINATED DOSE-RESPONSE

We map the dose-response curve of partial history dropping by sweeping two complementary protocols: **SELECTIVE-DROP** drops the **EARLIEST** K fraction of prior assistant messages, and **DROP-LATEST** drops the **LATEST** K fraction (Table 2). Three observations emerge.

First, neither curve is a smooth concave dose-response. Dropping any partial fraction (0.25 to 0.75) yields gap recovery between 4% and 23%, far short of the full removal at $K = 1.0$ which recovers 87%. The cliff at $K = 1.0$ is the dominant feature.

Second, the **EARLIEST**-vs-**LATEST** asymmetry is small but consistent: at $K = 0.5$, both yield 0.235 vs 0.242; at $K = 0.75$, **EARLIEST** yields 0.333 while **LATEST** yields 0.278. Removing earliest commitments is slightly more effective, supporting a weak “earliest-as-root-anchor” interpretation. But the magnitude is small (≈ 5 pp), and even at $K = 0.75$, retained later commits still anchor downstream—no partial removal escapes the chain.

Third, the implication is sharper than a graded dose-response: *any* retained partial-info commitment dominates attention at decision time, regardless of position. Interventions that keep “most recent” or “attention-important” commits (common heuristics in dialogue compression and KV eviction literature) will fail. Only complete removal of the assistant-history chain recovers performance.

Fourth, the cliff replicates cross-architecture. On Llama-3.1-8B (Table 2 bottom block), partial drops at $K=0.5$ recover 9.7% (EARLIEST: 0.185) and 1.3% (LATEST: 0.145) of the gap to Llama Concat (0.615), neither significantly above Sharded (Drop-EARLIEST $\Delta=+0.014$, $p=0.401$ ns; Drop-LATEST $\Delta=-0.007$, $p=0.541$ ns; bootstrap details in Appendix B). Drop-EARLIEST vs Drop-LATEST is itself null on Llama ($\Delta=+0.031$, $p=0.261$ ns), strengthening the binary-cliff claim: the ineffectiveness of partial drops is not Qwen-specific.

Per-turn evidence. Tracking accuracy turn-by-turn within the SHARDED protocol (Appendix F, Figure 2, $n=80$ math+data2text conversations) shows a non-monotone trajectory: turn 0–1 are very low (0.080, 0.115; insufficient evidence), turn 2 rises sharply to 0.430 as more shards have been revealed, then turn 3 *dips back* to 0.321 and the curve plateaus around 0.32–0.40. The dip is the per-turn signature of accumulated commitment bias: as more prior responses enter the context, their cumulative likelihood factor drags the posterior *away* from correctness even as the user reveals additional evidence. Cross-turn sentence ablation on $n=20$ failed conversations identifies specific sentences whose individual removal flips the final answer wrong-to-right (mean flip rate 0.04; 4/20 conversations have detectable single-sentence anchors), consistent with the cliff: any retained commitment can perpetuate contamination.

Why is the dose-response a cliff? A formal account.

Proposition 1 (Commitment-Contamination Cliff). *Let y be the correct answer, $\mathcal{E}$ the full user-revealed evidence at decision time, and $a_{1:T}$ the sequence of prior assistant responses, each generated under a strict evidence subset $\mathcal{E}_{<t} \subsetneq \mathcal{E}$. Under standard chat rendering, the LLM posterior decomposes as*

$$p(y \mid \mathcal{E}, a_{1:T}) \propto \underbrace{p(y \mid \mathcal{E})}_{\text{evidence-conditioned}} \cdot \underbrace{\prod_{t=1}^T p(a_t \mid y, \mathcal{E}_{<t})}_{\text{commitment likelihood}}.$$

For any retained subset $S \subseteq \{1, \dots, T\}$ with $|S| > 0$, the term $\prod_{t \in S} p(a_t \mid y, \mathcal{E}_{<t})$ persists and biases $p(y \mid \mathcal{E}, a_S)$ whenever any a_t committed to $y' \neq y$. The unique uncontaminated fixed point is $S = \emptyset$ —full assistant-history removal.

Remark 1 (Textual reweighting cannot truncate). *Appending an “this response is outdated” annotation to a_t changes the token distribution of the factor $p(a_t \mid y, \mathcal{E}_{<t})$ but leaves it non-zero. The Marked-History null result (Section 4.3) is the empirical confirmation of this prediction.*

The cliff at $K = 1.0$ is therefore not an empirical surprise: Proposition 1 predicts it as the unique bias-free operating point. Any retained assistant turn—regardless of position or recency—keeps the commitment likelihood active. This is why common heuristics such as “keep only the most recent k turns” or “keep attention-important tokens” fail: they leave at least one commitment factor in the posterior. Partial truncation interpolates between contaminated states; only $S = \emptyset$ exits.

5.4 PARETO: MPCHR VS ACCURACY ACROSS PROTOCOLS

Table 3 reports the three cache metrics alongside accuracy for all protocols (Qwen2.5-14B unless noted), and Figure 1 visualizes the (mPCHR, Accuracy) trade-off. **Recap attains the highest mPCHR among multi-turn protocols (0.851) but is dominated on accuracy (0.625):** retaining assistant history is the sole factor distinguishing it from Fresh-Last, which reaches 0.771. CACHE-GUARD attains accuracy 0.875 with mPCHR 0.858 and VCHR 0.993, comparable to FRESH-EVERY (0.875 / 1.000 / 1.000) but with structured eligibility that retains valid artifacts when present (a property F1 cannot exercise—we discuss the limitation in Section 7). MARKED-HISTORY achieves mPCHR comparable to Recap (0.755) but accuracy stays at 0.307—a text-level marker on stale commitments cannot override their attention-anchoring effect, ruling out the format-confusion hypothesis (H3) in favor of the attention-anchoring hypothesis (H1).

5.5 THE SALIENCE TRAP: ATTENTION-IMPORTANT TOKENS ARE THE HARMFUL ONES

An attention probe on Qwen2.5-7B (failed sharded BFCL conversations, $n = 22$ after OOM filtering of an initial $n = 30$ sample, HuggingFace eager attention, averaged across heads

Table 2: Selective-drop dose-response on multi.heldout. K is the fraction of prior assistant messages dropped. Drop-EARLIEST removes earliest K , Drop-LATEST removes latest K . $K = 0$ is the standard Sharded protocol; $K = 1.0$ Drop-EARLIEST is Fresh-Last. The cliff at $K = 1.0$ dominates and replicates on Llama-3.1-8B (bottom block). Llama block uses $K \in \{0, 0.5, 1.0\}$ (three-point sweep); Qwen block uses four intermediate values.

K	Drop-EARLIEST	Drop-LATEST	Symmetric
<i>Qwen2.5-14B</i>			
0.00 (Sharded)	0.208	0.208	—
0.25	0.354	0.253	EARLIEST >
0.50	0.235	0.242	tied
0.75	0.333	0.278	EARLIEST >
1.00 (Fresh-Last)	0.771	0.771	—
<i>Llama-3.1-8B (cross-arch replication)</i>			
0.00 (Sharded)	0.139	0.139	—
0.50	0.185	0.145	tied (ns)
1.00 (Fresh-Last)	0.651	0.651	—




Figure 1: (mPCHR, Accuracy) trade-off across protocols. Red markers are baselines and ablations (all dominated). Blue/green markers (FRESH-LAST, FRESH-EVERY, CACHEGUARD, 7B FRESH-LAST) sit on the Pareto frontier (dashed). RECAP attains the highest mPCHR among multi-turn baselines but is dominated on accuracy. The standard SHARDED protocol’s high mPCHR coincides with the worst accuracy: cache reuse and task correctness are fundamentally separate. Note: 7B FRESH-LAST uses Qwen2.5-7B as the assistant model; all other points use Qwen2.5-14B.

and layers) finds that prior assistant provisional commitments receive 8.5% more attention density per token than user evidence (0.00141 ± 0.00019 vs. 0.00130 ± 0.00014 , mean $\pm$ SE), and an attention density indistinguishable from the system prompt (0.00140 ± 0.00004). A naive comparison would expect user-evidence to dominate: it is what the model is supposed to condition on. Instead the model spends the same per-token attention on its own past commitments as on either of the two authoritative sources. The over-attended tokens here are content-anchored (provisional answers across the conversation) rather than position-anchored, distinguishing the signature from position-zero attention sinks (Xiao et al., 2024; Sun et al., 2025).

This directly indicts attention-salience-based KV compression: H2O (Zhang et al., 2023), StreamingLLM (Xiao et al., 2024), and SnapKV (Li, 2024) all rank tokens by attention contribution and preferentially retain high-attention tokens. In multi-turn agent settings these are precisely the harmful anchors, not the valid evidence. Validity requires an epistemic source label, not an attention score—the same conclusion reached independently by Wang & Sun (2025), who show prompt-level

Table 3: Three cache metrics and accuracy across protocols on multi_heldout (Qwen2.5-14B except 7B rows). PCHR is a message-level proxy for vLLM’s token-level prefix-cache hit rate. CACHEGUARD ties FRESH-EVERY on LiC and retains a Pareto advantage through typed eligibility (cf. BFCL in Section 5.7 and Mistral/Gemma in Table 5). Drop-K variants omitted; see Table 2.

Protocol	mPCHR	VCHR	Acc	Comment
Sharded	0.751	0.606	0.208	dominated
Concat	0.000	1.000	0.854	1-turn ceiling
Recap	0.851	0.842	0.625	high mPCHR, dominated
Marked-History	0.755	0.570	0.307	H1 over H3
Fresh-Last	0.788	0.930	0.771	frontier
Fresh-Every	1.000	1.000	0.875	frontier
7B Fresh-Last	0.758	0.914	0.917	cap. inversion
CacheGuard	0.858	0.993	0.875	frontier, structured

instructions to “ignore earlier inputs” fail to dislodge proactive interference. Full measurements are in Appendix E.

5.6 CROSS-BENCHMARK: SYNTHETIC BENCHMARKS (F1 AND F2)

To rule out LiC-specific artifacts, we constructed F1-additive (30 instances: multi-turn integer constraints across 4–6 turns, Sharded floor 0.609, Concat ceiling 0.977) and F2-corrective (Sharded 0.767, Concat 1.000). Every stateless protocol (FRESH-LAST, FRESH-EVERY, CACHEGUARD) reaches 1.000 on both, confirming the LiC finding generalizes and that removing *all* provisional commitments is the critical factor. Full tables are in Appendix A.

5.7 CROSS-BENCHMARK: STATEFUL TOOL USE (BFCL v4)

We evaluate on BFCL v4 `multi_turn_base` (Yan et al., 2024), a stateful function-calling benchmark in which a model issues API calls against mock backends (GorillaFileSystem, TwitterAPI, etc.) across up to five turns; correctness is judged against the final mock-API state via `multi_turn_checker`. Unlike LiC, the ground-truth object is external state mutated by tool calls, not an assistant-generated answer.

Setup. We implement a `ProtocolQwenHandler` adapter (subclassing `QwenFCHandler` from `bfcl-eval`) that tags each message with its block type (USER, ASST_TOOL_CALL, ASST_TEXT, TOOL) and turn index. At query time the adapter renders only prior turns according to the protocol rule, while always preserving the current turn’s assistant+tool exchanges intact (required for the inner tool-call step loop to converge). Qwen2.5-14B-Instruct is served via vLLM in function-calling mode; we report results on all $N = 200$ instances.

Why tool results are irreducible state. In BFCL, a `role=tool` message records the verified output of a prior API call—the outcome of `mv`, `grep`, or a state-mutating call is ground truth from the mock backend. The user never re-states tool results in subsequent turns. Any protocol that discards these messages at render time (FRESH-EVERY replaces all prior history with a concatenated user-only prompt; FRESH-LAST does the same at the final turn) renders the prior API state irrecoverable to the model.

CACHEGUARD’s eligibility controller classifies `role=tool` messages as always-valid (source: verified external state, not a provisional agent claim) and retains them across turns. It drops `asst_text` rationale from prior turns, which carries the anchoring risk identified in LiC but is not required for tool-state tracking. SHARDED retains everything including rationale.

Results. Table 4 shows the results. Both purging protocols are highly significantly worse than SHARDED (paired bootstrap, $p < 0.001$); CACHEGUARD matches SHARDED ($\Delta = -0.015$, $p = 0.66$, ns) and exceeds FRESH-EVERY by 19.0 pp ($p < 0.001$). The result holds on the long-context subset (CACHEGUARD= 0.200, SHARDED= 0.200, FRESH-EVERY= 0.067, $N = 30$).

Table 4: BFCL v4 `multi_turn_base` ($N = 200$, Qwen2.5-14B-Instruct). TOOL-RESULT-ONLY (TRO) retains user messages and `role=tool` results but drops all assistant messages. The ordering TRO>FE and CG>TRO isolates the contribution of (i) tool results and (ii) paired tool-call records.

Protocol	Acc	vs. Sharded
Sharded	0.295	—
Fresh-Last	0.120	−0.175***
Fresh-Every	0.090	−0.205***
Tool-Result-Only	0.210	−0.085**
CacheGuard	0.280	−0.015 (ns)

Ablation: tool-call records vs. tool results. TOOL-RESULT-ONLY (TRO) retains user messages and `role=tool` result messages but drops all assistant messages—including the `role=assistant tool_call` records that log which API was invoked with which arguments. TRO reaches 0.210: well above FRESH-EVERY (0.090, $\Delta = +0.120$, $p < 0.001$) but significantly below both SHARDED (0.295, $\Delta = -0.085$, $p = 0.004$) and CACHEGUARD (0.280, $\Delta = -0.070$, $p = 0.007$). This three-level ordering decomposes tool-state validity: tool results carry substantial information (TRO $\gg$ FE), but the *paired* tool-call record is also needed (CG > TRO). A tool result such as “FileNotFound” is only unambiguous when paired with the call that produced it. CACHEGUARD’s eligibility rule—retain `role=tool` *and* the preceding `role=assistant tool_call`, drop `asst_text rationale`—captures exactly the irreducible state.

Token cost and turn-depth degradation. FRESH-EVERY’s accuracy collapse is not compensated by any input-token saving: its mean input context (58,231 tokens) is 7.3% larger than SHARDED (54,254), while its mean output length nearly doubles (1,510 vs. 706 tokens), reflecting repeated failed tool-call attempts. CACHEGUARD is only 4.9% larger than SHARDED on input tokens and 25% larger on output, yet ties on accuracy. The penalty from discarding tool state grows with conversation depth: at four turns FRESH-EVERY reaches 0.02 (vs. SHARDED 0.34, CACHEGUARD 0.28); at five or six turns it reaches 0.00. SHARDED and CACHEGUARD hold at 0.21–0.34 across all turn-depth buckets.

Two-regime picture. The LiC and BFCL results define two distinct validity regimes:

- **Stateless-reasoning regime (LiC):** prior assistant outputs are provisional partial-info claims; retaining them anchors later responses. Optimal: FRESH-EVERY (0.875). CACHEGUARD $\approx$ FRESH-LAST (both drop unverified claims).
- **Stateful-tool regime (BFCL):** prior tool results are verified external state; discarding them is catastrophic. Optimal: SHARDED or CACHEGUARD (0.295/0.280, ns difference). FRESH-EVERY drops to 0.090.

No fixed-retention protocol is Pareto-optimal across both regimes. CACHEGUARD’s typed eligibility—retain verified external state, drop unverified agent claims—is the only protocol that achieves competitive accuracy in both.

5.8 CROSS-MODEL GENERALIZATION

The strongest cross-architecture result is on Llama-3.1-8B: SHARDED drops to 0.139 against a Concat ceiling of 0.615; FRESH-EVERY reaches 0.777, exceeding the ceiling by 0.156 ($p = 0.004$, **robust to best-of-4 control at $\Delta = +0.125$**). We test across model size (Qwen-7B vs. 14B), architecture family (Llama-3.1-8B), and three additional open-weight families (Mistral-7B, Gemma-2-9B), with the user simulator (Qwen2.5-7B) held fixed throughout. On Qwen-7B, FRESH-LAST reaches 0.917 and FRESH-EVERY 0.896—both nominally exceed the same-model Concat ceiling (0.771); a multi-attempt control (Appendix C) shows this exceedance largely reduces to best-of-N advantage, so we report 7B as “stateless protocols make multi-attempt usable” rather than “stateless extracts more than Concat”. On Llama FRESH-LAST (0.651) reaches the Concat ceiling (0.615; ns) and FRESH-EVERY significantly exceeds it. Cross-architecture paired bootstrap (Appendix B, $n = 23$ –24) gives Fresh-Last > Sharded ($p < 0.001$) and Fresh-Every > Fresh-Last ($p = 0.009$).

CACHEGUARD also transfers cleanly across architectures. On Llama-8B it reaches 0.704, beating Sharded ($\Delta=+0.533$, $p < 0.001$) and the Concat ceiling ($\Delta=+0.080$, $p = 0.032$), and sits between Fresh-Last and Fresh-Every on the Pareto frontier ($FE > CG$, $\Delta=+0.076$, $p = 0.043$)—the same relative position it occupies on Qwen-14B (Section 5.4). On Mistral-7B CACHEGUARD reaches 0.544, exceeding FRESH-EVERY ($\Delta=+0.015$); on Gemma-2-9B it reaches 0.741, marginally exceeding FRESH-EVERY ($\Delta=+0.008$). On these two models typed eligibility outperforms uniform forgetting—suggesting that some assistant-block content is genuinely valid evidence and the typed renderer correctly retains it. The four architecture families (Qwen, Llama, Mistral, Gemma) giving consistent CacheGuard placement on the Pareto frontier rule out model-family-specific quirks as the source of either the anchoring effect or the validity-aware mitigation.

Capability inversion. A counterintuitive consequence of the anchoring hypothesis is that under SHARDED, a stronger model’s lift from forgetting can *exceed* that of a weaker one—larger models generate more coherent partial-info commitments, which become stronger anchors. The clearest cross-architecture case is Llama-3.1-8B FRESH-EVERY 0.777 vs. same-model Concat ceiling 0.615: a +0.162 exceedance robust to best-of-4 control (+0.125). Llama’s SHARDED floor (0.139) is also lower than Qwen-14B’s (0.208) despite Llama-8B’s strong single-turn ceiling (0.615), implying the multi-turn protocol penalizes Llama more, not less. The implication for practitioners: scaling up the base model does not, on its own, escape the anchoring effect; it can amplify it. The protocol fix is therefore not redundant with model improvement.

Two additional families confirm the pattern (Table 5). Mistral-7B-Instruct-v0.2: Sharded floor 0.148, FRESH-EVERY 0.529 exceeds Concat 0.396 ($\Delta=+0.133$). Gemma-2-9B-it: Sharded floor 0.169, FRESH-LAST 0.746 and FRESH-EVERY 0.733 both exceed Concat 0.656 ($\Delta=+0.090$ and $+0.077$). Across five models in four architecture families (Qwen-7B/14B, Llama-3.1-8B, Mistral-7B, Gemma-2-9B), every non-Qwen-14B model has a stateless protocol that meets or exceeds its same-model Concat ceiling.

5.9 STATISTICAL SIGNIFICANCE

We use paired bootstrap ($B = 10,000$, common-task pairing) on the 24 multi-heldout instances. All paper-level claims pass at $\alpha = 0.05$: CacheGuard > Sharded ($p < 0.001$); Fresh-Last > Marked-History ($p < 0.001$); CacheGuard > Recap ($p = 0.005$); Fresh-Every > Recap ($p = 0.005$); 7B Fresh-Last > 14B Fresh-Last ($p = 0.019$). Cross-architecture: Llama-8B Fresh-Last > Sharded ($\Delta=+0.482$, $p < 0.001$); Fresh-Every > Fresh-Last ($\Delta=+0.125$, $p = 0.009$); CacheGuard > Sharded ($\Delta=+0.533$, $p < 0.001$); CacheGuard > Concat ceiling ($\Delta=+0.080$, $p = 0.032$). Cross-architecture mechanism: Drop-EARLIEST $K=0.5$ vs Sharded ns ($p = 0.401$, cliff replicates); Marked-History vs Sharded ns ($p = 0.402$, H1 replicates). The Marked-History vs. Sharded null ($p = 0.34$) and Drop-EARLIEST vs. LATEST null at $K = 0.5$ ($p = 0.60$) confirm that the cliff at $K = 1.0$ is the dominant feature, not slope or format. Full bootstrap table with confidence intervals is in Appendix B.

6 DISCUSSION

Validity is a separate cache axis from physical reuse and attention salience. Existing serving systems (Kwon et al., 2023; Zheng et al., 2024; Liu, 2025) treat all cached prefix tokens as equally reusable; KV compression methods (Zhang et al., 2023; Xiao et al., 2024; Li, 2024) preferentially retain attention-important tokens. Our finding implies that, in multi-turn agentic settings, attention-important tokens may be precisely the harmful anchors. The selective-drop curve (Section 5.3) makes this concrete: dropping high-attention tokens (which arrive at $K = 0.5/0.75$) does not help because those tokens are themselves products of an anchored chain. The right axis to optimize is not physical reusability or attention salience, but *epistemic validity under the current task state*—a property orthogonal to both.

The binary-cliff finding rules out compromise mitigations. A graded dose-response would have invited compromise: keep half the assistant history to balance cache reuse and quality. Our data show this compromise is empty. Partial drops at $K = 0.25$ to 0.75 all yield gap-recovery between 4% and 23%, far below the cliff at $K = 1.0$ (87%). The implication is that production

Table 5: Cross-model generalization. Qwen-14B rows from multi_heldout ($N=2$); all other rows from multi_mca24 ($N=4$). User simulator is Qwen2.5-7B throughout; the actions grader is Qwen2.5-7B for Qwen-7B and Qwen2.5-14B otherwise. Stateless protocols meet or exceed the same-model Concat ceiling on every non-Qwen-14B model; CACHEGUARD matches or exceeds FRESH-EVERY on Mistral and Gemma. [†]Mistral-7B Math Concat (0.344) falls marginally below Sharded (0.367) on small- N noise; the aggregate ordering (Concat avg 0.396 > Sharded avg 0.148) holds.

Protocol	Math	Code	Actions	Avg
14B Sharded (reference)	0.500	0.000	0.125	0.208
14B Concat (reference)	0.938	0.625	1.000	0.854
14B Fresh-Last	0.875	0.438	1.000	0.771
14B Fresh-Every	0.938	0.688	1.000	0.875
7B Sharded	0.500	0.103	0.000	0.201
7B Concat (ceiling)	0.781	0.531	1.000	0.771
7B Fresh-Last	1.000	0.750	1.000	0.917
7B Fresh-Every	0.938	0.750	1.000	0.896
Llama-8B Sharded	0.318	0.000	0.100	0.139
Llama-8B Concat (ceiling)	0.594	0.250	1.000	0.615
Llama-8B Fresh-Last	0.742	0.241	0.969	0.651
Llama-8B CacheGuard	0.667	0.444	1.000	0.704
Llama-8B Fresh-Every	0.815	0.517	1.000	0.777
Mistral-7B Sharded	0.367	0.038	0.038	0.148
Mistral-7B Concat (ceiling)	0.344 [†]	0.031	0.812	0.396
Mistral-7B Fresh-Last	0.481	0.048	0.800	0.443
Mistral-7B Fresh-Every	0.500	0.154	0.933	0.529
Mistral-7B CacheGuard	0.500	0.350	0.781	0.544
Gemma-2-9B Sharded	0.375	0.100	0.033	0.169
Gemma-2-9B Concat (ceiling)	0.719	0.281	0.969	0.656
Gemma-2-9B Fresh-Last	0.839	0.400	1.000	0.746
Gemma-2-9B Fresh-Every	0.786	0.414	1.000	0.733
Gemma-2-9B CacheGuard	0.759	0.464	1.000	0.741

dialogue managers which summarize and compress history—common in long-context applications and agentic memory frameworks—will continue to suffer the anchoring effect as long as any answer fragment from a partial-info turn remains. A correct deployable mitigation must remove *all* prior assistant content at decision time, or replace it with structured representations that the model does not interpret as authoritative evidence.

Implications for production multi-turn LLM systems. Our finding differentiates by conversation structure: compositional turns (each decision depends on accumulating evidence) suffer maximally from anchoring, while local self-contained turns do not (Huang et al., 2026). Production multi-turn LLM systems—chatbots, coding assistants, virtual agents in customer service or medical triage—may be paying a substantial accuracy tax for a default they could change in a single client-side rendering patch. Our protocols are deployable today and require no model retraining.

Multi-turn dialogue is sometimes overhead, not a feature. A stronger reading of our results: many “multi-turn capability” benchmark scores measure how well a model tolerates an unnecessary protocol overhead, not how well it reasons across turns. FRESH-EVERY matches or exceeds the same-model Concat ceiling on every studied model, with the exceedance robust to a multi-attempt control on Llama-8B (Appendix C)—in this regime, multi-turn structure is at best inert and often harmful. The optimal action for a chat-protocol system facing an underspecified multi-turn input is often not to reason multi-turn, but to detect that the multi-turn structure is redundant and collapse it to a single equivalent query. This reframes “multi-turn capability” itself: the desideratum is not turn-by-turn reasoning improvement, but better detection of when forgetting (and collapse)

is appropriate. A non-trivial fraction of reported multi-turn benchmark drops are therefore protocol artifacts—measurable, but addressable without model changes. Reproducible multi-turn comparisons should specify the context-rendering policy alongside the model.

Active forgetting as a missing primitive. Generation has been the focus of LLM research; retention (KV cache reuse) is now a serving-systems focus; *active forgetting* is the missing third operation. Standard chat protocols provide no mechanism to selectively drop content other than full-history truncation, and KV compression methods that approximate forgetting via attention salience choose precisely the wrong tokens (Section 5.5). Our cliff finding rules out gradient-based forgetting heuristics; what works is a typed, content-aware drop policy. Production deployments and serving stacks should expose forgetting policies as configurable, queryable properties—a VCHR/HCHR pair or analog should be reportable alongside PCHR. CACHEGUARD is one instantiation; learned controllers, dialogue-summarization variants, and explicit memory-architecture wrappers are others. The unifying claim: retention is no longer the only relevant memory operation.

Connections to working memory, databases, and operating systems. The forgetting primitive has analogs across disciplines that sharpen its status as a first-class operation. In cognitive science, *directed forgetting* (Bjork, 1972) and ACT-R activation decay (Anderson & Lebière, 1998; Kosinski, 2025) establish that biological working memory implements explicit eviction; Wang & Sun (2025) show LLMs lack this and exhibit log-linear retrieval decay under proactive interference, independently corroborating our anchoring effect. In databases, multi-version concurrency control without garbage collection causes unbounded version-chain growth and read degradation; transactional logs distinguish *provisional* (pre-commit) from *durable* (post-commit) state—a distinction directly mirrored in our `assistant_claim` vs. `tool_observation` typing. In operating systems, virtual memory requires both a paging layer (cache reuse) and an eviction policy (forgetting); MemGPT (Packer et al., 2023) and MemOS (Li et al., 2025b) adopt this framing for LLM context management but stop short of typed eviction. We argue active forgetting completes the analogy: chat-protocol multi-turn is broken without it, in the same sense that an OS without a page-eviction policy is broken—not as a missing optimization, but as a missing structural component.

7 LIMITATIONS

Open-weight only; two benchmarks. Our cross-model evaluation covers five open-weight models (Qwen2.5-7B, Qwen2.5-14B, Llama-3.1-8B, Mistral-7B-Instruct-v0.2, Gemma-2-9B-it); generalization to closed-source families (GPT, Claude) is plausible but unverified. Primary results use LiC and BFCL v4, supplemented by F1/F2 synthetics. Cross-benchmark validation on τ^2 -Bench (Barrès, 2025), MultiChallenge (Weston, 2025), and additional BFCL subsets is left to future work. The miss-param subset ($N = 30$) shows no CacheGuard advantage (0.167 vs. Sharded 0.233), consistent with our prediction that CacheGuard targets tool-state retention, not function-spec robustness.

Reasoning models. We do not fully evaluate reasoning models (o3, DeepSeek-R1, QwQ). Laban et al. (2025) report that on LiC these models degrade similarly to non-reasoning models, with extra test-time compute providing no recovery—consistent with our framing: the reasoning chain is generated *after* the anchored prior commitments are rendered into the context window, so anchors are in-context during the entire trace. The prediction is that FRESH-LAST should still help, since the anchoring mechanism is upstream of the reasoning trace. Preliminary evaluation of DeepSeek-R1-Distill-7B on `multi_mca24` is underway (vLLM on GPU port 8006); full cross-model results including reasoning models are left to follow-up.

Long-context models. A natural objection is that 128K+ context models trivially absorb the multi-turn conversation. Wang & Sun (2025) show otherwise: proactive interference produces log-linear retrieval-accuracy decay even when the target value is adjacent to the query and well within context. This decoupling between context-length and effective-retrieval is consistent with our anchoring effect. Direct verification on Gemini 1.5/Llama-3.1-128K is left to follow-up.

CacheGuard eligibility boundary on mixed-API tasks. Among the $N = 200$ BFCL base instances, CACHEGUARD and SHARDED disagree on 33 instances: SHARDED-only correct (18)

and CACHEGUARD-only correct (15). The SHARDED-only cases are concentrated in combined MessageAPI+TradingBot tasks, where the assistant’s prior *text* rationale appears to carry portfolio or thread-context summaries that the mock backend does not expose via explicit tool results. Our current eligibility rule (drop all `asst_text`) therefore discards valid informational state in these cases. The CACHEGUARD-only cases are predominantly single-API TradingBot tasks, where dropped rationale was a stale incorrect intermediate plan. This suggests that `asst_text` blocks occupy a spectrum between pure provisional claims and partially-valid state summaries; the binary drop rule is an approximation that wins on average but leaves room for finer-grained eligibility based on content type.

Approximate PCHR, unpopulated HCHR, and correlational attention. Our mPCHR is a message-level approximation to vLLM’s token-level prefix cache hit rate; the two diverge for FRESH-EVERY where user content extends rather than rebuilds per turn. HCHR, as defined in Section 4.1, requires counterfactual per-block ablations to evaluate whether removing a block flips the outcome; this computation is expensive at scale. A five-conversation pilot on failed math sharded conversations (Qwen2.5-14B, $n = 5$) finds mean HCHR = 0.16: individual block drops rarely recover accuracy, consistent with Proposition 1’s prediction that any retained commitment perpetuates contamination (partial removal fails; full removal succeeds). Full-scale empirical HCHR measurement is left for follow-up. The attention measurements in Section 5.5 ($n = 22$ failed sharded BFCL conversations) are correlational; causal interventions (attention patching, KV-block ablation) would more rigorously establish the anchoring mechanism but are left for follow-up.

Rule-based eligibility controller. CACHEGUARD’s eligibility logic is hand-written and could fail on non-standard interaction patterns. A learned eligibility classifier is a natural extension but out of scope.

8 CONCLUSION

We argued that multi-turn LLM degradation is most parsimoniously explained not as a capability deficit, but as the absence of an *active forgetting* mechanism in standard chat protocols: agents persist their own provisional commitments as context with the same standing as user evidence. Three forgetting metrics (PCHR, VCHR, HCHR) characterize what is retained vs. should be dropped; a cliff-shaped selective-drop curve characterizes the mechanism (forgetting is binary, not graded); a five-line forgetting policy (FRESH-LAST) closes 87% of the multi-turn gap on Lost-in-Conversation, and FRESH-EVERY matches or exceeds the same-model single-turn ceiling. The forgetting effect and CACHEGUARD both replicate across five models and four architecture families (Qwen-7B/14B, Llama-3.1-8B, Mistral-7B, Gemma-2-9B); on Mistral and Gemma CACHEGUARD exceeds FRESH-EVERY, evidence that typed eligibility can outperform uniform forgetting. On stateful tool-use (BFCL) the same controller reverses cleanly: FRESH-EVERY collapses while CACHEGUARD matches the strongest baseline, demonstrating that typing is structurally necessary, not optional. Two implications follow: (i) multi-turn dialogue is sometimes overhead rather than a feature, and many multi-turn benchmark drops measure protocol weakness rather than model weakness; (ii) chat protocols and serving stacks should expose forgetting as a first-class primitive alongside generation and retention.

REFERENCES

- et al. Agrawal. Gepa: Reflective prompt evolution can outperform reinforcement learning. *arXiv preprint arXiv:2507.19457*, 2025. Accepted at ICLR 2026.
- John R. Anderson and Christian Lebière. *The atomic components of thought*. Lawrence Erlbaum, 1998.
- Anonymous. Contextual drag: How errors in the context affect LLM generation. In *Under review at ICLR*, 2026. OpenReview zpiYsPVDIV.
- et al. Barres. τ^2 -bench: Evaluating conversational agents in a dual-control environment. *arXiv preprint arXiv:2506.07982*, 2025.

- Robert A. Bjork. Theoretical implications of directed forgetting. *Coding processes in human memory*, pp. 217–235, 1972.
- ByteDance-Tsinghua-SIA. MemAgent: Reshaping long-context LLM with multi-conv RL-based memory agent. In *ICLR*, 2026. Oral; arXiv:2507.02259.
- et al. Chen. Learning to clarify: Multi-turn conversations with action-based contrastive self-training. In *ICLR*, 2025.
- Hao Chen et al. AI psychosis: Persistent incorrect beliefs in long-horizon LLM agents. *arXiv preprint arXiv:2604.13860*, 2026.
- Jenny Y. Huang, Leshem Choshen, Ramon Astudillo, Tamara Broderick, and Jacob Andreas. Do LLMs benefit from their own words? *arXiv preprint arXiv:2602.24287*, 2026.
- Omar Khattab et al. A survey of context engineering for large language models. *arXiv preprint arXiv:2507.13334*, 2025.
- Sungwoo Kim et al. Causal token attribution in language models via gradient tracing. *arXiv preprint arXiv:2510.02334*, 2025.
- et al. Kosinski. Human-like remembering and forgetting in LLM agents: An ACT-R-inspired memory architecture. In *HAI*, 2025. ACM DL: 10.1145/3765766.3765803.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention. In *SOSP*, 2023.
- Philippe Laban, Hiroaki Hayashi, Yingbo Zhou, and Jennifer Neville. Llms get lost in multi-turn conversation. *arXiv preprint arXiv:2505.06120*, 2025.
- Yoonho Lee, Roshen Nair, Qizheng Zhang, Kangwook Lee, Omar Khattab, and Chelsea Finn. Meta-harness: End-to-end optimization of model harnesses. *arXiv preprint arXiv:2603.28052*, 2026.
- et al. Li. SnapKV: LLM knows what you are looking for before generation. In *NeurIPS*, 2024.
- et al. Li. Mitigating conversational inertia in multi-turn agents through context bias calibration. *arXiv preprint arXiv:2602.03664*, 2026a.
- et al. Li. Intent mismatch causes llms to get lost in multi-turn conversation. *arXiv preprint arXiv:2602.07338*, 2026b.
- Jinghao Li et al. Thought anchors: Which part of the reasoning chain matters most? *arXiv preprint arXiv:2506.19143*, 2025a.
- Zhiyu Li et al. MemOS: An operating system for memory-augmented generation (MAG) in large language models. *arXiv preprint arXiv:2505.22101*, 2025b.
- et al. Liu. LMCACHE: An efficient KV cache layer for enterprise-scale LLM inference. *arXiv preprint arXiv:2510.09665*, 2025.
- Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12, 2024. arXiv:2307.03172.
- Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. MemGPT: Towards LLMs as operating systems. In *NeurIPS*, 2023. arXiv:2310.08560.
- Joon Sung Park, Joseph C. O’Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *UIST*, 2023. arXiv:2304.03442.
- Yichuan Sun et al. When attention sink emerges in language models: An empirical view. In *ICLR*, 2025. arXiv:2410.10781.

- Amos Tversky and Daniel Kahneman. Judgment under uncertainty: Heuristics and biases. *Science*, 185(4157):1124–1131, 1974.
- Chupei Wang and Jiaqiu Vince Sun. Unable to forget: Proactive interference reveals working memory limits in LLMs beyond context length. *arXiv preprint arXiv:2506.08184*, 2025. ICML Workshop on Long-Context Foundation Models.
- et al. Weston. MultiChallenge: A realistic multi-turn conversation evaluation benchmark challenging to frontier LLMs. In *ACL Findings*, 2025. *arXiv:2501.17399*.
- Chien-Sheng Wu, Richard Socher, and Caiming Xiong. Transferable multi-domain state generator for task-oriented dialogue systems. In *ACL*, 2019.
- Qingyun Wu et al. Recursively summarizing enables long-term dialogue memory in large language models. *arXiv preprint arXiv:2308.15022*, 2023. Published in Neurocomputing 2025.
- Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. Efficient streaming language models with attention sinks. In *ICLR*, 2024.
- et al. Xu. Learning to forget: Sleep-inspired memory consolidation for resolving proactive interference in large language models. *arXiv preprint arXiv:2603.14517*, 2026.
- Fanjia Yan, Huanzhi Mao, Charlie Cheng-Jie Ji, Tianjun Zhang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. Berkeley Function Calling Leaderboard. https://gorilla.cs.berkeley.edu/blogs/8_berkeley_function_calling_leaderboard.html, 2024.
- et al. Zhang. Aflow: Automating agentic workflow generation. In *ICLR*, 2025. Oral.
- et al. Zhang. Breaking contextual inertia: Reinforcement learning with single-turn anchors for stable multi-turn interaction. *arXiv preprint arXiv:2603.04783*, 2026.
- Zhenyu Zhang, Ying Sheng, Tianyi Zhou, Tianlong Chen, Lianmin Zheng, Ruisi Cai, Zhao Song, Yuandong Tian, Christopher Ré, Clark Barrett, Zhangyang Wang, and Beidi Chen. H2O: Heavy-hitter oracle for efficient generative inference of large language models. In *NeurIPS*, 2023.
- Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Sun, Jeff Huang, Cody Hao Yu, Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett, and Ying Sheng. SGLang: Efficient execution of structured language model programs. *arXiv preprint arXiv:2312.07104*, 2024.

A SYNTHETIC BENCHMARK RESULTS (F1 AND F2)

F1-additive poses a multi-turn underspecified math task: across 4–6 turns, the user reveals constraints on a positive integer (divisibility, parity, bounds, digit-sum); the assistant must commit to the smallest integer satisfying all revealed constraints. Every assistant turn before the final shard is, by construction, a partial-info commitment. Concat ceiling is 0.977 and the Sharded floor is 0.609—a 37 pp multi-turn drop that mirrors the LiC pattern. F2-corrective requires the assistant to track explicit user retractions (“Actually, the limit is X, not Y”), creating an even simpler anchoring scenario where the final correct constraint directly contradicts the partial one. Both benchmarks confirm that removing *all* provisional commitments fully recovers accuracy.

On both F1 and F2 every stateless protocol attains 1.000. This does not differentiate CACHEGUARD from FRESH-EVERY, because neither benchmark generates an assistant artifact that the user later references back. The BFCL regime (Section 5.7) provides the stress-test where tool results represent irreducible verified state that must be retained.

B PAIRED BOOTSTRAP SIGNIFICANCE TESTS

All bootstrap tests use $B = 10,000$ iterations with common-task pairing on the 24 multi_heldout instances (Qwen2.5-14B unless noted). p is one-sided.

Table 6: F1 additive-constraints and F2 corrective-constraints synthetic benchmarks (Qwen2.5-14B, $N = 2$ sims/instance). F1: user reveals additive constraints across 4–6 turns. F2: user gives an initial bound then explicitly retracts and corrects it. Both reproduce a multi-turn drop fully closed by any stateless protocol.

Protocol	F1 (additive)		F2 (corrective)	
	Acc	Gap recovery	Acc	Gap recovery
Concat (ceiling)	0.977	100%	1.000	100%
Sharded (floor)	0.609	0%	0.767	0%
Fresh-Last	1.000	106%	1.000	100%
Fresh-Every	1.000	106%	1.000	100%
CacheGuard	1.000	106%	1.000	100%

Table 7: Paired bootstrap pairwise tests on multi_heldout (Qwen2.5-14B unless noted). p is one-sided. Significance at $\alpha = 0.05$.

Comparison	Δ	95% CI on Δ	p	sig
CacheGuard > Sharded	+0.645	[+0.476, +0.810]	< 0.001	***
Fresh-Last > Marked-History (H1>H3)	+0.416	[+0.229, +0.604]	< 0.001	***
CacheGuard > Recap	+0.250	[+0.042, +0.458]	0.005	**
Fresh-Every > Recap	+0.250	[+0.042, +0.458]	0.005	**
7B Fresh-Last > 14B Fresh-Last (cap. inv.)	+0.146	[+0.021, +0.292]	0.019	*
<i>Within-Qwen-7B (cross-size capability inversion, $n = 24$)</i>				
7B Fresh-Last > 7B Sharded	+0.719	[+0.562, +0.865]	< 0.001	***
7B Fresh-Every > 7B Sharded	+0.698	[+0.531, +0.854]	< 0.001	***
7B Fresh-Last > 7B Concat (ceiling)	+0.146	[+0.052, +0.260]	0.008	**
7B Fresh-Every > 7B Concat (ceiling)	+0.125	[+0.052, +0.219]	0.007	**
<i>Cross-architecture (Llama-3.1-8B-Instruct, $n = 23$ common tasks)</i>				
Llama-8B Fresh-Last > Sharded	+0.482	[+0.333, +0.630]	< 0.001	***
Llama-8B Fresh-Every > Sharded	+0.612	[+0.449, +0.761]	< 0.001	***
Llama-8B Fresh-Every > Fresh-Last	+0.125	[+0.042, +0.219]	0.009	**
Llama-8B Fresh-Every > Concat (ceiling)	+0.156	[+0.062, +0.260]	0.004	**
Llama-8B CacheGuard > Sharded	+0.533	[+0.370, +0.692]	< 0.001	***
Llama-8B CacheGuard > Concat (ceiling)	+0.080	[+0.010, +0.160]	0.032	*
Llama-8B Fresh-Every > CacheGuard	+0.076	[+0.000, +0.160]	0.043	*
<i>Cross-architecture mechanism (Llama-3.1-8B, $n = 22-24$)</i>				
Llama-8B Fresh-Last > Drop-EARLIEST K=0.5	+0.479	[+0.323, +0.635]	< 0.001	***
Llama-8B Fresh-Last > Marked-History (H1>H3)	+0.507	[+0.348, +0.663]	< 0.001	***
Marked-History vs Sharded (14B)	+0.049	[−0.095, +0.214]	0.340	ns
Llama-8B Marked-History vs Sharded	+0.008	[−0.064, +0.102]	0.402	ns
Llama-8B Drop-EARLIEST K=0.5 vs Sharded	+0.014	[−0.101, +0.130]	0.401	ns
Llama-8B Drop-LATEST K=0.5 vs Sharded	−0.007	[−0.167, +0.156]	0.541	ns
Llama-8B Drop-EARLIEST vs LATEST K=0.5	+0.031	[−0.052, +0.125]	0.261	ns
Llama-8B Fresh-Last vs Concat	+0.031	[−0.052, +0.115]	0.266	ns
Llama-8B CacheGuard vs Fresh-Last	+0.049	[−0.035, +0.146]	0.163	ns
Drop-EARLIEST-50 vs Drop-LATEST-50	0.000	[−0.091, +0.091]	0.601	ns

C MULTI-ATTEMPT CONTROL: BEST-OF-N RESCORING

A possible confound for “stateless protocol exceeds Concat ceiling” is multi-attempt: Sharded/Fresh-Last/Fresh-Every all expose K turns of verifier checks within a single run, whereas Concat is single-shot. Across $N = 4$ independent reps per task, we rescore each protocol two ways: (i) *mean* of per-rep accuracy (default), (ii) *best-of- N* (task succeeds if any of the N reps succeeds).

Table 8: Best-of- N rescoring on multi_heldout ($N = 4$ reps per task). Anchoring (Sharded $\ll$ stateless) survives the control on both models; “Fresh-Last/Fresh-Every $>$ Concat” partially reduces to multi-attempt advantage on Qwen-7B but persists on Llama-8B.

Protocol	Qwen-7B		Llama-3.1-8B	
	mean	best-of-4	mean	best-of-4
Sharded	0.198	0.333	0.157	0.274
Concat	0.771	0.875	0.615	0.750
Fresh-Last	0.917	0.917	0.646	0.875
Fresh-Every	0.896	0.917	0.771	0.875
CacheGuard	—	—	0.694	0.792

Two takeaways: (i) the anchoring claim (Sharded $\ll$ stateless protocols) is robust to the multi-attempt control: even best-of-4 Sharded (0.274 / 0.333) is far below best-of-4 stateless (0.792–0.917). (ii) The narrower “stateless $>$ same-model Concat” claim is partially explained by multi-attempt on Qwen-7B (best-of-4 Concat 0.875 nearly closes the gap to FL/FE 0.917) but persists on Llama-8B (best-of-4 FE 0.875 vs. Concat 0.750, $\Delta=+0.125$). We thus restrict the strong “exceeds same-model ceiling” claim to Llama-8B, and frame the 7B finding as “stateless makes multi-attempt usable” rather than “stateless extracts more information than Concat”.

D MARKED-HISTORY ABLATION

MARKED-HISTORY keeps every prior assistant message in context but prefixes each with the literal string [OUTDATED PARTIAL-INFO COMMITMENT]. If gap recovery were comparable to FRESH-LAST, format/instruction-following would be the dominant mechanism (H3). On multi_heldout (Qwen2.5-14B, $n = 37$ tasks), Marked-History reaches 0.307—only 4.9 pp above Sharded (0.208) and 46.4 pp below Fresh-Last (0.771). The pairwise paired-bootstrap test against Fresh-Last is highly significant ($p < 0.001$, $\Delta = +0.416$) while the test against Sharded is not significant ($p = 0.34$, $\Delta = +0.049$). A textual marker is insufficient to override the underlying anchoring effect, supporting H1 (attention-mediated) over H3 (format-confusion).

Cross-architecture replication on Llama-3.1-8B. The same test on Llama-3.1-8B yields Marked-History 0.170, only 3.1 pp above Llama Sharded (0.139) and 48.1 pp below Llama Fresh-Last (0.651). Bootstrap: Marked-History vs Sharded $\Delta=+0.008$, $p=0.402$ ns; Fresh-Last $>$ Marked-History $\Delta=+0.507$, $p < 0.001$. The H1 conclusion replicates: a textual marker fails on a distinct architecture family, and is dominated by full forgetting by an effect size comparable to the 14B result.

Positive control: [FACT] injection in user turns. The Marked-History null could be attributed to general instruction-following difficulty rather than an asymmetry specific to assistant-turn content. We test this by injecting a *user-turn* hint into the same failed conversations: the final user shard is appended with [FACT: The correct answer is: {gold}]. Use this in your final answer.]

We select $n = 35$ conversations where Sharded failed (math $n = 7$, data2text $n = 16$, actions $n = 12$; Qwen2.5-14B) and run three conditions per conversation: original Sharded, Marked-History ([OUTDATED] on all asst turns), and FACT injection (gold in last user turn). Scoring follows the per-task canonical evaluator (exact-match for math/actions, RougeL for data2text).

Condition	Math	Data2Text	Actions
Original Sharded	0.143	0.306	0.000
[OUTDATED] on asst turns	0.143	0.301	0.000
[FACT] in last user turn	0.571	0.548	0.167

Across all three task types, [OUTDATED] leaves accuracy effectively unchanged from the Sharded baseline (0.143 vs. 0.143, 0.301 vs. 0.306, 0.000 vs. 0.000), while [FACT] yields consistent gains

(+43 pp math, +25 pp data2text, +17 pp actions). The model *can* read and apply user-turn instructions; the failure is specific to annotating assistant-content. Remark 1 stated this prediction from Proposition 1—the FACT control is its empirical confirmation.

E ATTENTION PROBE RESULTS

We load Qwen2.5-7B-Instruct via HuggingFace transformers (attn_implementation=eager) on $n = 22$ failed sharded BFCL function-calling conversations (initial sample $n = 30$, with 8 OOM-filtered for length). We forward each conversation’s full last-turn input and read attention weights from the final input position to all prior tokens, averaging across heads and layers, then bucket by block type. Conversations are processed independently; we report per-conversation block-type means aggregated across conversations.

Table 9: Attention density per block type (Qwen2.5-7B, $n = 22$ failed sharded BFCL conversations; mean $\pm$ standard error). Prior assistant tokens receive attention density indistinguishable from system tokens and 8.5% above user evidence—small in absolute magnitude but qualitatively at the wrong end of the salience spectrum given that user shards are the authoritative evidence source. Attention-salience-based KV compression heuristics that retain high-attention tokens would therefore preferentially keep the harmful anchors.

Block type	Avg attention density per token
System prompt	0.00140 ± 0.00004
Assistant provisional commits	0.00141 ± 0.00019
User evidence (shards)	0.00130 ± 0.00014

The earlier pilot ($n = 5$) reported a 52% gap between asst-commits and user-evidence; the scale-up to $n = 22$ shrinks this to 8.5%, indicating the original pilot over-estimated the magnitude due to small-sample variance in user-shard length. The qualitative ordering—asst-commits $\geq$ system $>$ user-evidence—is robust across the larger sample. Causal interventions (attention patching that zeros out attention to prior assistant tokens, then measuring whether Sharded performance recovers toward Fresh-Last) would more rigorously establish the mechanism and are left for follow-up.

F COMMITMENT ANCHOR IDENTIFICATION

F.1 METHODOLOGY: CROSS-TURN SENTENCE-LEVEL ABLATION

A sentence s in assistant turn a_t ($t < T$) is a *commitment anchor* if its removal from the conversation history causes the model’s final-turn response to flip from incorrect (score < 0.5) to correct (score ≥ 0.5). This is a causal, counterfactual criterion applied cross-turn.

Procedure. For each failed sharded conversation: (i) extract all non-final assistant turns $a_1, \dots, a_{T-1}$; (ii) split each turn into sentences using punctuation-boundary and newline splitting (min. 15 chars); (iii) for each sentence $s_{t,i}$, construct a modified conversation $\hat{M}^{(t,i)}$ with $s_{t,i}$ removed; (iv) re-run the model on $\hat{M}^{(t,i)}$ and re-grade via the same LLM grader; (v) mark $s_{t,i}$ as an anchor iff $\text{score}(\hat{M}^{(t,i)}) \geq 0.5$ while $\text{score}(M) < 0.5$. The *sentence flip rate* is the fraction of tested sentences that are anchors per conversation. This extends the Thought Anchors methodology (Li et al., 2025a) from intra-turn CoT (their setting) to cross-turn assistant history (our setting); the causal criterion and sentence-splitting logic are shared.

The key difference: Li et al. (2025a) remove sentences from a single reasoning trace and test whether the final answer changes; we remove sentences from *earlier turns* and test whether the model’s response on the *next/final turn* changes. The cross-turn gap is larger—the removed sentence may have occurred 3–5 turns earlier, yet still causally determine the outcome, consistent with Proposition 1.

F.2 PER-TURN ACCURACY TRAJECTORY

Figure 2 reports per-turn accuracy on $n=80$ SHARDED conversations (math and data2text, Qwen2.5-14B-Instruct graded by Qwen2.5-7B-Instruct). Three features are diagnostic. First, accuracy at turns 0 and 1 is very low (0.080, 0.115)—the model commits to an answer with insufficient evidence and that commitment is wrong. Second, turn 2 shows a sharp rise to 0.430—more shards have arrived and the additional evidence pushes a substantial fraction of conversations to a correct response. Third, turn 3 dips back to 0.321 and the curve plateaus around 0.32–0.40 thereafter despite further user evidence: the cumulative weight of prior commitments now drags accuracy *down* even as more information becomes available. This non-monotone trajectory—peak at turn 2, dip at turn 3, plateau below the same-model Concat ceiling—is the per-turn signature of the cliff predicted by Proposition 1: each retained commitment factor adds bias, and at some point the bias outweighs the marginal informational gain from a new shard. Within-conversation, 67/80 (84%) reach a correct response at some turn but only 13/80 retain it as the final response (Section 5.2); the remaining 54 commit again and lose the answer.




Figure 2: Per-turn accuracy on $n=80$ SHARDED conversations (math and data2text, Qwen2.5-14B-Instruct, n at each turn is the count of conversations reaching that turn). Turn 0–1 is the underspecified-commitment regime. Turn 2 marks a sharp rise as enough shards have been revealed for a correct attempt. Turn 3 dips back and the curve plateaus around 0.32–0.40 despite more user evidence—accumulated commitments anchor the model below the Concat ceiling. Turns 6–7 have small n (15, 3); the rise there is dominated by easy-task selection bias.

F.3 AGGREGATE FINDINGS

We ran the cross-turn sentence ablation on $n=20$ failed sharded conversations (math and data2text only, conversations capped at ≤ 8 assistant turns; up to 30 randomly sampled sentences per conversation; Qwen2.5-14B-Instruct). The mean per-conversation flip rate is 0.04: 4% of tested sentences are commitment anchors. 4 of 20 conversations (20%) contain at least one anchor sentence. The distribution is heavy-tailed: one data2text conversation accounts for 17 of 23 anchor sentences (a single-turn entry whose every fabricated date is a flip-causing commitment), while three other conversations contribute one to three anchors each. The 80%/20% split between conversations with no detectable single-sentence flip and conversations with at least one is not evidence against the cliff: Proposition 1 predicts that any non-empty S persists contamination, so even a single retained commitment can anchor the outcome—most failed conversations have a small number of unique anchor sentences, not none.

F.4 CASE STUDY: CLARIFICATION-LOOP ANCHOR IN MATH

The clearest single-anchor flip in the experiment comes from a math conversation (sharded-GSM8K/651, 7 assistant turns, Qwen2.5-14B-Instruct). The first user shard underspecifies the problem (asks about totals without revealing the input numbers); the assistant correctly recognizes this and asks for clarification, but the form of the clarification commits the model to a stance (“data is missing”) that persists through the rest of the conversation:

Turn 1 (Asst, original SHARDED): “To accurately calculate Ariadne’s total sales for March and April, we need specific data regarding her sales for those months. Since the actual sales figures are not provided in your question, I will outline the steps to calculate her total sales once you have the necessary data. [...] steps follow...”

Turn 7 (Asst, SHARDED final): response repeats the steps abstractly without committing to a numeric answer. Score: 0 (gold = “5”).

Turn 7 (Asst, after dropping either underlined sentence): model now produces a numeric answer. Score: 1.

Anchor sentences (both flip the outcome individually):

- “To accurately calculate Ariadne’s total sales for March and April, we need specific data regarding her sales for those months.” (Turn 1, sentence 0)
 - “Since the actual sales figures are not provided in your question, I will outline the steps to calculate her total sales once you have the necessary data.” (Turn 1, sentence 1)
-

Both sentences are early commitments to a meta-stance (“need more data” / “will outline the steps”). When either sentence is in context, every subsequent turn defaults back to outlining steps rather than computing; when removed, the same model on the same conversation produces the correct numeric answer. The anchor here is structural rather than factual—the model is anchored on a *response mode*, not a wrong value. This generalizes the F1-additive intuition: commitments come in many forms (numerical, factual, modal), and any of them can chain forward.

Companion data2text case (fabrication anchor). A second conversation (sharded-totto-Pinckney) shows a complementary failure mode where the model fabricates specific factual claims early and then anchors. The gold output records that John M. Pinckney died on April 24, 1905; the model committed to “passed away on February 16, 1905” in a turn-4 response and produced fabricated successor-dates downstream. 17 of 30 tested sentences in this conversation flip wrong-to-right when removed—most are fabricated dates, names, or speculative process descriptions. This case illustrates that on knowledge-intensive tasks, anchoring extends to manufactured evidence: the model treats its own earlier fabrications as authoritative input on later turns.

G CACHEGUARD ALGORITHM

Algorithm 1 CACHEGUARD renderer (called before every assistant turn)

Require: Conversation trace T , current turn index t

Ensure: Messages M to be passed to the LLM

```

1: Parse  $T$  into typed blocks  $B_1, B_2, \dots, B_n$  via Block Parser
2: Annotate contradicted_by and referenced_by relations across blocks
3:  $M \leftarrow [B_{\text{system}}]$ 
4: for each block  $B_i$  with  $B_i.\text{turn} < t$  do
5:   if  $B_i.\text{source} \in \{\text{user}, \text{tool}\}$  and not contradicted then
6:      $M \leftarrow M \parallel B_i$ 
7:   else if  $B_i.\text{type} = \text{artifact}$  and  $B_i.\text{referenced\_by} \neq \emptyset$  then
8:      $M \leftarrow M \parallel B_i$ 
9:   else
10:    hide  $B_i$  ▷ provisional commitments to side-buffer
11:   end if
12: end for
13:  $M \leftarrow M \parallel B_{\text{current user turn}}$ 
14: return  $M$ 

```
